

Lab 06 - Deploying your model

Lab 06 - Deploying your model

The last lab was a bit of a tangent, but it was to show some of the quirks of working with Docker.

Now that we have a good grasp on the basics, let's do the same for our model.

Step 1

In the last set of labs, we wrapped our model up in a Flask script. Now we have an API we can ping our model and get a response. What we don't want to do now is have to run the model each time manually. Let's put it in a docker container.

Have a go at doing what you think you need to base on the last lab.

Remember the general steps are:

1. Create a dockerfile
2. Create the application code
3. Create a requirements.txt file
4. Build the image
5. Check that the model is running

If you do get stuck here is the answer:

Lab 06 - Deploying your model

You cheat! I am just kidding. It is complicated.

Let's walk through the solution end-to-end.

Step 1

1. Create the following files:

```
Dockerfile
DiabetesModelFlask.py
requirements.txt
```

2. In the Dockerfile we will want to add the following:

```
FROM python:3-onbuild
COPY . usr/src/app
CMD ["python", "DiabetesModelFlask.py"]
```

3. In DiabetesModelFlask.py we want to add the following:

```
# Product Service
from flask import Flask, request, jsonify
from flask_restful import Resource, Api
import _pickle as pickle
import numpy as np
from sklearn import datasets, linear_model

app = Flask(__name__)
api = Api(app)

PickleModelPath = './regression.pkl'

@app.route('/score')
def apicall():
    var1 = float(request.args.get('var1'))
    with open(PickleModelPath, 'rb') as k:
        PickleModel = pickle.load(k)
    Answer = PickleModel.predict(var1)
    return jsonify(Answer.tolist())

@app.route('/train')
def apicalled():
    # Load the diabetes dataset
    diabetes = datasets.load_diabetes()
    # Use only one feature
    diabetes_X = diabetes.data[:, np.newaxis, 2]
    # Split the data into training/testing sets
    diabetes_X_train = diabetes_X[:-20]
    diabetes_X_test = diabetes_X[-20:]
    # Split the targets into training/testing sets
    diabetes_y_train = diabetes.target[:-20]
    diabetes_y_test = diabetes.target[-20:]
```

Lab 06 - Deploying your model

```
# Create linear regression object
regr = linear_model.LinearRegression()
# Train the model using the training sets
regr.fit(diabetes_X_train, diabetes_y_train)
# Make predictions using the testing set
diabetes_y_pred = regr.predict(diabetes_X_test)
with open(PickleModelPath, 'wb') as f:
    pickle.dump(regr, f)
return "Model has been retrained. Run /score to score model"

if __name__ == '__main__':
    app.run(host='0.0.0.0', port=5071, debug=True)
```

Now our requirements are slightly more difficult. We have a lot more modules which need to be added.

4. Add the following to requirements.txt

```
pandas==0.18.1
scikit-learn==0.19.1
numpy==1.15.4
scipy==0.18.1
Flask==0.12
flask_restful==0.3.5
```

Step 2

5. Navigate to the location of your docker file.

For me that looks like this

```
G:
cd G:\GitHubProjects\workshop-
ModelManagement\MachineLearningFromModelToProduction\Docker\DiabetesProductionDock
er
```

6. Build our model in to an image

```
docker build -t productiondiabetes .
```

This will build our image.

7. Run our image as a container

```
docker run -d --name productiondiabetes -p 5071:5071 productiondiabetes
```

8. Navigate to localhost:5071

You will see an error. That is because part of our web service is missing. If you train the model first, you will be able to execute the model.

Lab 06 - Deploying your model

9. Navigate to the following url

```
localhost:5071/train?var1=1  
localhost:5071/score?var1=1
```

Press ctrl+c to end the session.

Lab done.